Module 2: Building the Database

Designing relational databases, tables and constraints

Module Objectives

Understand how relational databases are structured in SQL Server

Configure database files with appropriate size and growth settings

Define tables with suitable data types and indexing

Safeguard data integrity using keys and constraints

Constructing a Relational Database

- A database groups related tables, views and other objects
- Data is stored in data files and tracked in transaction log files
- File size and growth settings affect performance and manageability
- Schemas organise objects and help manage security and ownership

Choosing File Size and Growth Parameters

- Set an initial file size large enough to avoid frequent autogrowth
- Use fixed growth increments (MB) rather than percentage for larger databases
- Configure separate data and log files with appropriate locations
- Monitor file usage and adjust settings as the database grows over time

Working with Database Schemas

- A schema is a container for tables, views and other objects
- Schemaname.objectname identifies objects uniquely within a database
- Schemas can represent applications, modules or business areas
- Permissions can be granted at schema level to simplify security

Defining Tables

- Tables store rows of data with a fixed set of columns
- Each column has a name, data type and optional default value
- Primary key columns uniquely identify each row
- Design tables to reduce redundancy while keeping queries manageable

Specifying Appropriate Data Types

- Choose the smallest data type that can hold required values
- Use numeric types for quantities, money types for financial amounts
- Use date and time types instead of storing dates as strings
- Avoid using NVARCHAR(MAX) or similar large types unless necessary

Indexing Tables with Management Studio

- Indexes help SQL Server find rows quickly without scanning entire tables
- Create clustered indexes on key columns used for range searches and ordering
- Create nonclustered indexes on frequently filtered or joined columns
- Use SSMS to design, create and inspect indexes graphically

Adding, Modifying and Removing Columns

- Use ALTER TABLE to add new columns as requirements change
- Be careful when changing data types on existing columns
- Removing columns can break applications and reports that depend on them
- Plan schema changes and, if possible, test them in a non-production environment

Safeguarding Data with Constraints

- Primary key constraints enforce uniqueness of key values
- Foreign key constraints protect relationships between tables
- Check constraints validate data values against business rules
- Constraints prevent invalid data from being stored in the first place